


Prediction Systems

Dan Crankshaw

UCB RISE Lab Seminar

10/3/2015

Learning



Big Model

Timescale: minutes to days


Systems: offline and batch optimized

*Heavily studied ... major focus of the **AMPLab***

Learning




Training



Inference

Query

Decision




Application

Learning



Training

Inference



Application

Timescale: ~20 milliseconds


Systems: *online* and *latency* optimized
Less studied ...

Learning



Training


Feedback



Inference

Query

Decision




Application

Learning




Training



Inference

Decision



Application

Timescale: hours to weeks

Systems: combination of systems

Less studied ...

Feedback

Learning

Inference

Adaptive
(~1 seconds)

Responsive
(~10ms)

Feedback




Training




Query

Decision

Application



Big Model



Prediction Serving Challenges

- Complexity of deploying new models
 - New applications or products (**$0 \rightarrow 1$ models**).
 - New data, features, model family: (**$N \rightarrow N+1$ models**).
 - *Why is it hard:* Frameworks not designed for low-latency serving, frameworks have **different APIs, different resource requirements, and different costs**.
- System Performance
 - Need to ensure **low-latency predictions, scalable throughput**. Deploying a new model can't degrade system performance.
- Model or Statistical Performance
 - Model Selection: Which models to use?
 - When to deploy a new model?
 - How to adapt to feedback?
 - *At a meta-level: what are the right metrics for measuring model performance?*

LASER: A Scalable Response Prediction Platform for Online Advertising

Agarwal et al. 2014



LASER Overview

- Top-down system design enforced by company organizational structure
- Picked a model (logistic regression) and built the system based on that choice
- Force data-scientists to use this model, express features in specialized configuration language
- Result: **System and model family are tightly coupled**

$$p_{ijt} = \frac{1}{1 + \exp(-s_{ijt})}$$

$$s_{ijt} = \omega + s_{ijt}^{1,c} + s_{ijt}^{2,c} + s_{ijt}^{2,\omega}$$

Addressing Deployment Complexity

- **Fixed Model Choice:** Can be hardcoded into system, no need for API to specify model
- **Configuration language:** specify feature construction in JSON-based configuration language
 - Restricts feature transformations to be built from component library
 - Allows for changes in pipeline without service restarts or code modification
 - Allows easy re-use of common features across an organization
 - Similar to PMML, PFA
- **Language details**
 - *Source*: translate data to numeric feature vectors
 - *Transformer*: Vector-to-vector transformations (transform, aggregate)
 - *Assembler*: Concatenates all feature pipelines together into single vector

Addressing System Performance

- **Precompute second-order interaction terms**
 - The LASER logistic regression model includes second order interaction terms between user and campaign features:

$$s_{ijt}^{2,c} = x_i' A c_j + \dots$$


- **Don't wait for delayed features**
 - Features can be delayed by slow DB lookup, expensive computation
 - *Solution: Substitute expected value for missing features and degrade accuracy, not latency*
 - *Solution: Cache precomputed scalar products in PRC, save overhead of re-computing features and dot products which are lazily evaluated*

Addressing Model Performance

- Decompose model into slowly-changing and quickly-changing components
 - Fast retraining of warm-start (quickly-changing) component of model without cost of full retraining

$$s_{ijt} = \boxed{\omega + s_{ijt}^{1,c} + s_{ijt}^{2,c}} + \boxed{s_{ijt}^{2,\omega}}$$

**Cold Start
Trained Offline** **Warm Start
Trained Online**

- Explore/Exploit with Thompson Sampling
 - Sometimes serve ads with low empirical mean but high-variance
 - Draw sample from posterior distribution over parameters and use sample to predict CTR instead of mode
 - In practice, hold Θ_c fixed and sample from Θ_w

Some Takeaways from LASER

- System performance is paramount in the broader application context
 - Slow page load has much larger impact on revenue than poor ad-recommendation
- AUC/accuracy is not always the most useful model performance metric
- The more assumptions you can make about your tools (software, models) the more tricks you can play (config language, shared features, warm-start/cold-start decomposition)
 - Safe for LASER to make these assumptions because they are enforced through extra-technological methods
 - Similar to some of the design choices we saw in Borg last week

Clipper

A Low-Latency Online Prediction Serving System

Daniel Crankshaw,

Xin Wang

Giulio Zhou

Michael Franklin,

Joseph E. Gonzalez

Ion Stoica



Goals of Clipper

- **Design Choice:** *General purpose, easy to use* prediction serving system
 - Generalize to many *different ML applications* (contrast to LASER which was designed to address LinkedIn's ad-targeting needs)
 - Generalize to *many frameworks/tools* for a single application
 - Don't tie the hands of data scientists developing models
 - Make it simple for a *data-scientist* to deploy a new model into production
- Given these design choices, maximize system and model performance using *model-agnostic* techniques

Clipper Generalizes Models Across ML Frameworks

Fraud
Detection



Content
Rec.



Personal
Asst.



Robotic
Control



Machine
Translation



Clipper

theano

Dato



KeystoneML

Caffe

TensorFlow

scikit
learn

dmlc
mxnet

VW

KALDI

Clipper Architecture

Applications



Predict

RPC/REST Interface

Observe

Clipper

theano

Dato



KeystoneML

Caffe

TensorFlow

scikit
learn


dmlc
mxnet



VW
KALDI

Clipper Architecture

Applications



Predict

RPC/REST Interface

Observe

Clipper



RPC

RPC

RPC



RPC

Model Wrapper (MW)




...

Clipper Architecture




Clipper Architecture



Caching

Adaptive Batching


Model Abstraction Layer



Approximate Caching

Adaptive Batching

Model Abstraction Layer




Common Interface → Simplifies Deployment:

- Evaluate models using original code & systems
- Models run in separate processes (Docker containers)
 - Resource isolation

Caching

Adaptive Batching

Model Abstraction Layer




Common Interface → Simplifies Deployment:

- Evaluate models using original code & systems
- Models run in separate processes
 - Resource isolation
 - Scale-out

Problem: frameworks optimized for **batch processing** not **latency**


Adaptive Batching to Improve Throughput

- Why batching helps:



A single page load may generate many queries

Hardware Acceleration



Helps amortize system overhead

- Optimal batch depends on:


- hardware configuration
- model and framework
- system load

Clipper Solution:


be as **slow** as **allowed**...

- Inc. batch size *until the latency objective is exceeded* (**Additive Increase**)
- If latency exceeds SLO cut batch size by a fraction (**Multiplicative Decrease**)

Adaptive Batching to Improve Throughput



Clipper Architecture



Selection Policy

Model Selection Layer


Goal:

*Maximize **accuracy** through **bandits** and **ensembles**,
online learning, and **personalization***




Incorporate feedback in real-time to achieve:

- **robust predictions** by combining multiple models & frameworks
- **online learning** and **personalization** by selecting and personalizing **predictions** in response to feedback

Learning



Slow Changing
Model




Feedback

Slow

Inference


Clipper

Fast Model
Selection per-
User



Application

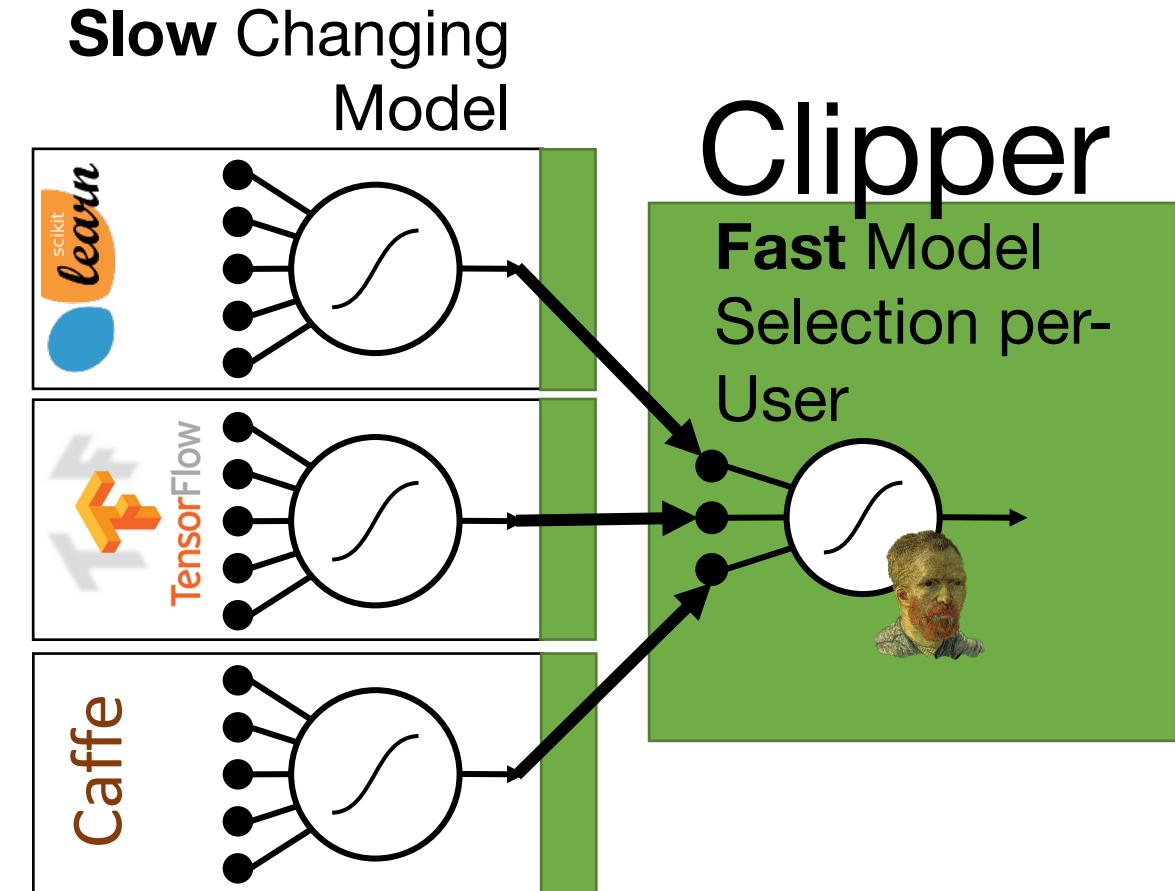
Fast Feedback



Model Selection Policy

Improves prediction **accuracy** by:

- Incorporates real-time **feedback**
- Estimates **confidence** of predictions
- Determines how to **combine** multiple **predictions**
 - e.g., choose best, average, ...
- enables frameworks to **compete**




Cost of Ensembles

Increased Load

- *Solutions:*
 - **Caching and Batching**
 - **Model Selection** prioritizes frameworks for load-shedding

Stragglers

- e.g., framework fails to meet SLO
- *Solution: Anytime* predictions
 - Selection policy must select/combine from *available* predictions
 - e.g., built-in ensemble policy
substitutes expected value



Limitations of Clipper

- Clipper does not address offline model retraining
- By treating deployed models as black boxes, Clipper forgoes the opportunity to optimize prediction execution of the models themselves or share computation between models
- Only performs coarse-grained tradeoffs of accuracy, robustness, and performance.

TensorFlow Serving

- Recently released open-source prediction-serving system from Google
- Companion to TensorFlow deep-learning ML framework
- Easy to deploy **TensorFlow Models**
- System automatically manages the lifetime of deployed models
 - Watches for new versions, loads and transfers requests to new models automatically
- System does not address model performance, only system performance (through batching)

TensorFlow Serving Architecture

Applications



Predict
↔

RPC/REST Interface

TensorFlow-Serving

Prediction Batching

RETIRED

V1



TensorFlow

V2



TensorFlow

New model
version trained


V3



TensorFlow

TensorFlow Serving Architecture

Applications



Predict
↔

RPC/REST Interface

TensorFlow-Serving

Prediction Batching

RETIRED



Other Prediction-Serving Systems

➤ **Turi**



- Company co-founded by **Joey**, Carlos Guestrin, and others to serve predictions from models (primarily) trained in the GraphLab Create framework
- Not open-source
- Recently acquired by Apple

➤ **Oryx**



- Developed by Cloudera for serving Apache Spark Models
- Implementation of Lambda Architecture with Spark and Spark Streaming to incrementally maintain models
- Open source

➤ **PredictionIO**



- Open-source Apache Incubating project, the company behind the project was recently acquired by Salesforce
- Built on Apache Spark, Hbase, Spray, ElasticSearch